

HPC Lab Week 3 Report 2

Name :	Prateek Agrawal
Roll Number :	CED18I040
Programming Environment :	OpenMP
Problem Statement :	Vector Dot Product
Date :	26th August 2021

Systems Specifications :

CPU Name :	Intel(R) Core(TM) i7-8750H CPU @ 2.20GHz
CPU Type :	Intel Coffeelake processor
CPU Stepping :	10
Number of Sockets: :	1
Cores per Socket :	6
Threads per core :	2
L1 Cache size	32 kB
L2 Cache size	256 kB
L3 Cache size:	9 MB
RAM	32 GB

Serial Code:

```
/*
 * @Author: prateek
 * @Date: 2021-08-26 15:51:41
 * @Last Modified by: prateek
 * @Last Modified time: 2021-08-26 16:10:52
 */

#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include <omp.h>
#define VECTOR_LENGTH 10000
int main (int argc, char *argv[])
{

    double arr_1[VECTOR_LENGTH] ;
    double arr_2[VECTOR_LENGTH] ;
    double start, end;
    for (int i = 0; i < VECTOR_LENGTH; i++)
    {
        arr_1[i] = arr_2[i] = i * 1.0;
    }

    double sum = 0;

    start = omp_get_wtime();

    for (int k = 0; k < 1000000; k++)
    {
        for (int i = 0; i < VECTOR_LENGTH; i++)
        {
            sum += arr_1[i] * arr_2[i];
        }
    }

    end = omp_get_wtime();
    printf("%lf", end - start);

    return 0;

}
```

Parallel Code : (Reduction)

```
/*
 * @Author: prateek
 * @Date: 2021-08-26 15:37:42
 * @Last Modified by: prateek
 * @Last Modified time: 2021-08-26 16:13:58
 */

#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include <omp.h>
#define VECTOR_LENGTH 10000
int main (int argc, char *argv[])
{

    double arr_1[VECTOR_LENGTH] ;
    double arr_2[VECTOR_LENGTH] ;
    double start, end;
    for (int i = 0; i < VECTOR_LENGTH; i++)
    {
        arr_1[i] = arr_2[i] = i * 1.0;
    }

    double sum = 0;

    start = omp_get_wtime();

    #pragma omp parallel for reduction (+:sum)
    for (int k = 0; k < 1000000; k++)
    {

        for (int i = 0; i < VECTOR_LENGTH; i++)
        {
            sum += arr_1[i] * arr_2[i];
        }
    }

    end = omp_get_wtime();
    printf("%lf", end - start);

    return 0;
}
```

Parallel Code : (Reduction + Critical Section)

```
/*
 * @Author: prateek
 * @Date: 2021-08-26 16:20:56
 * @Last Modified by: prateek
 * @Last Modified time: 2021-08-26 16:28:16
 */

#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include <omp.h>
#define VECTOR_LENGTH 10000
int main (int argc, char *argv[])
{

    double arr_1[VECTOR_LENGTH] ;
    double arr_2[VECTOR_LENGTH] ;
    double start, end;
    for (int i = 0; i < VECTOR_LENGTH; i++)
    {
        arr_1[i] = arr_2[i] = i * 1.0;
    }
    int i;
    double sum = 0;
    double psum = 0;

    start = omp_get_wtime();

    #pragma omp parallel shared(arr_1,arr_2,sum) private(i,psum)
    {
        psum = 0.0;
        #pragma omp for
        for (int k = 0; k < 1000000; k++)
        {
            for ( i = 0; i < VECTOR_LENGTH; i++)
            {
                psum += arr_1[i] * arr_2[i];
            }
            #pragma omp critical
            {
                sum+=psum;
            }
        }
    }
}
```

```

    }
}
}
end = omp_get_wtime();
printf("%lf", end - start);

return 0;

}

```


Compilation and Execution:

For enabling OpenMP environment use -fopenmp flag while compiling using gcc.


```
g++ -fopenmp vector_dot_product_parallel_reduction.cpp
```

Table 1 - Parallel Code with Reduction

NUM Threads	Execution Time	Speed-Up	Parallelization Fraction
1	25.349	1	
2	14.586	1.73789935554641	84.9185372204032
4	7.718	3.2844001036538	92.7373860901811
6	6.81	3.7223201174743	87.7620418951438
8	6.062	4.18162322665787	86.9552476006379
10	5.586	4.53795202291443	86.6262530627989
12	4.951	5.11997576247223	87.7839900444343
16	5.831	4.3472817698508	82.130261548779
32	5.667	4.47308981824599	80.1487365411119
64	5.985	4.23542188805347	77.6021345195672
128	6.331	4.00394882325067	75.6154011262616
150	5.851	4.33242180823791	77.4344513014426



NUMBER OF THREADS vs SPEED-UP (Parallel with Reduction)




NUMBER OF THREADS vs Execution Time (Parallel with Reduction)

Inference: (Note: Execution time, graph and inference will be based on hardware configuration)


- At thread count 12 maximum speedup is observed.
- If thread count is more than 12 then the execution time increases/decreases slightly and tapers out.

Table 2 - Parallel Code with Reduction and Critical Section

Threads	Execution Time	Speed-Up	Parallelization_Fraction
1	25.538	1	
2	13.977	1.82714459469128	90.5395880648446
4	7.782	3.2816756617836	92.7036833999008
6	6.901	3.70062309810172	87.5730284282246
8	6.143	4.15725215692658	86.7950281373416
10	6.247	4.08804226028494	83.9315703831328
12	5.916	4.31676808654496	83.819477569967
16	6.391	3.99593177906431	79.9728509149764
32	6.53	3.91087289433384	76.8312369422922
64	6.459	3.95386282706301	75.8941235407678
128	6.28	4.06656050955414	76.0029673242838
150	6.145	4.15589910496339	76.4474679396042



NUMBER OF THREADS vs SPEED-UP (Parallel with Reduction and Critical Section)



NUMBER OF THREADS vs Execution Time (Parallel with Reduction and Critical Section)

Inference: (Note: Execution time, graph and inference will be based on hardware configuration)

- At thread count 12 maximum speedup is observed.
- If thread count is more than 12 then the execution time increases/decreases slightly and tapers out.